

CSCI 561

Foundation for Artificial Intelligence

Probabilistic Decision Making


Professor Wei-Min Shen
University of Southern California

Outline

- **Probabilistic decision making**
 - **Motivation**
 - **Decision Problems**
 - **Markov Decision Process (MDP)**
 - **Value Iteration**
 - **Policy Iteration**

Slides contributions from: Brian C. Williams
(MIT 16.410), Manuela Veloso,
Reid Simmons, & Tom Mitchell, CMU

How Might a Mouse Search a Maze for Cheese?




- State Space Search?
- As a Constraint Satisfaction Problem?
- Goal-directed Planning?
- Linear Programming?

What is missing?

Action and Sensor Models (review)


1. Actions
2. Percepts (observations)
3. States
4. Appearance: states \rightarrow observations
5. Transitions: (states, actions) \rightarrow states
6. Current State



•What about the goals?

Utility Value of States \Leftrightarrow Goal Information

1. Actions
2. Percepts (observations)
3. States
4. Appearance: states \rightarrow observations
5. Transitions: (states, actions) \rightarrow states
6. Current State
7. **Rewards**: $R(s)$ or $R(s,a)$ (related to the goals, given to the agent from env)
- (8) **Utility Value of States**: $U(s)$ (how good for the goals, must compute)

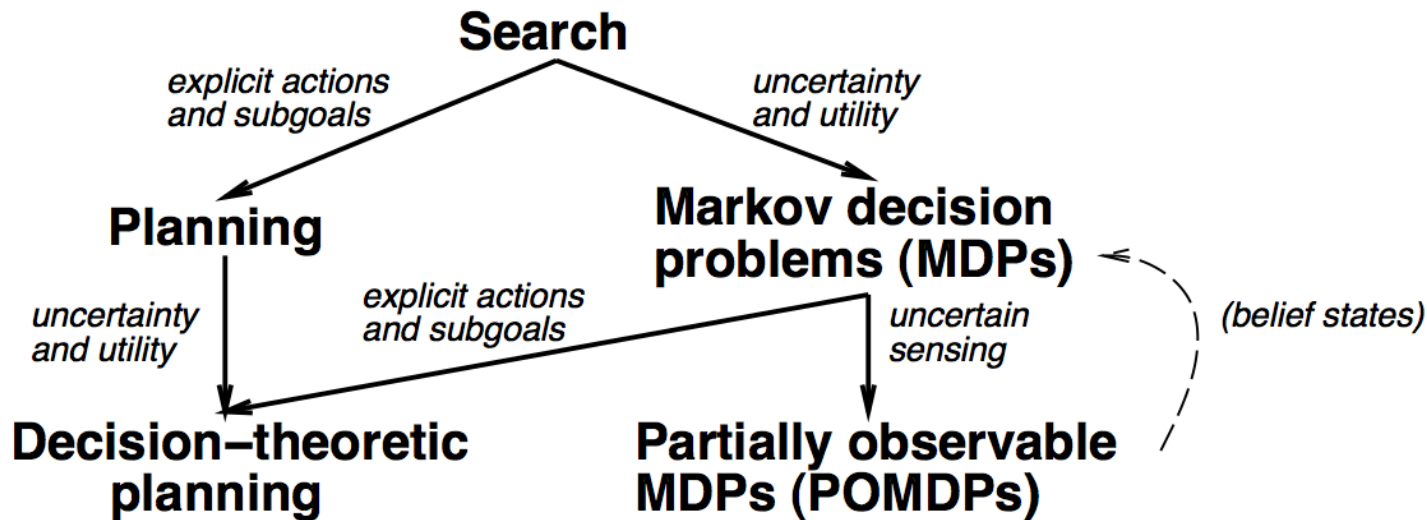


Note that rewards can be given from states $R(s)$ or state-action $R(s,a)$

The Key Representations

- Model the environment by **States, Actions, Percepts**, and
- **Transition model** $\phi = P(s_{t+1}|s_t, a_t)$ may be probabilistic
- **Sensor Model** $\theta = P(z|s)$ may be probabilistic
- States may have **Utility Value** $U(s)$ or $V(s)$
- Agents may receive **Reward** r (implicit “goals”)
 - When entering a state: $s \rightarrow r$
 - When performing an action in a state: $(s, a) \rightarrow r$
- Behaviors may be represented as **Policy**: $s \rightarrow a$
- Objective: find the optimal policy based on utilities or rewards

Sequential Decision Problems



Markov Decision Process (MDP)

- A MDP consists of
 - State S
 - Actions A
 - Transition Model $\Phi(s'|s,a)$
 - Initial State s_0 Probability Distribution π
 - ~~• Sensor Model $\theta(z|s)$~~ // a typical MDP has no sensor model
 - A reward function $R(s)$ or $R(s, a)$ // we use $R(s, a)$ in today's lecture


Maximum Expected Utility (MEU) and Rational Agents

- Every state has a utility value $U(s)$
- The expected utility of an action given the current evidence or observation e , is the average utility value of the outcomes, weighted by the probability that the outcome occurs:

$$EU(a | e) = \sum_{s'} P(result(a) = s' | a, e) U(s')$$

- The principle of maximum expected utility (MEU) is that a **rational** agent should choose the action that maximizes its expected utility:

$$action = \arg \max_a EU(a | e)$$




POMDP: Sensor Model, Belief States

- A Partially Observable MDP (POMDP) is defined as follows:
 - A set of states S (with an initial state s_0)
 - A set of Actions: $A(s)$ of actions in each state s
 - A transition model $P(s'|s,a)$, or $T(s,a,s')$
 - A reward function $R(s)$, or $R(s,a)$
 - A sensor model $P(e|s)$ ← ←
 - A belief of what the current state is $b(s)$ ← ←
- Belief States (where am I now? What is my current state?):
 - If $b(s)$ was the previous belief state, and the robot does action “ a ” and then perceives a new evidence “ e ”, then the new belief state:

$$b'(s') = \alpha P(e | s') \sum P(s' | s, a) b(s)$$

where α is a normalization constant making the belief states sum to 1

Goals, Rewards, Utilities, Policies



- Goals
 - Given to the agent from the problem statements
- Rewards
 - Given to the agent, designed based on the goals
- Utility values for states
 - Computed by the agent, based on the rewards
- Policies
 - Computed or learned by the agent
 - Used by the agent to select its actions
 - The better a policy, the more rewards it collects

Compute Utilities from Rewards over Time

- Utility Value = sum of all future rewards

- Add the rewards as they are

$$U_h = ([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

- Discount the far-away rewards in the future

$$U_h = ([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

- The expected future utility value $U^\pi(s)$ obtained by executing a policy π starting from s

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t) \right]$$

The Optimal Policy

$$\pi_s^* = \arg \max_{\pi} U^\pi(s)$$

Ideas in this lecture

- **Problem** is to **accumulate rewards**, rather than to achieve goal states.
- **Approach** is to generate **reactive policies** for how to act in all situations, rather than plans for a single starting situation.
- **Policies** fall out of **value functions**, which describe the greatest **lifetime reward** achievable at every state.
- **Value functions** are **iteratively approximated**.

MDP Examples: TD-Gammon [Tesauro, 1995]

Learning Through Reinforcement

Learns to play Backgammon

States:


- Board configurations (10^{20})

Actions:

- Moves


Rewards:

- +100 if win
 - - 100 if lose
 - 0 for all other states
-
- Trained by playing 1.5 million games against self.
-
- ➔ Currently, roughly equal to best human player.



MDP Examples: Aerial Robotics [Feron et al.]



Computing a Solution from a Continuous Model



Markov Decision Processes

- Motivation
- What are Markov Decision Processes (MDPs)?
 - Models
 - Lifetime Reward
 - Policies
- Computing Policies From a Model
- Summary

Example MDP



Model $M_{ij}^a \equiv P(j|i, a)$ = probability that doing a in i leads to j

Each state has a *reward* $R(i)$

= -0.04 (small penalty) for nonterminal states

= ± 1 for terminal states

Example MDP

In search problems, aim is to find an optimal *sequence*

In MDPs, aim is to find an optimal *policy*

i.e., best action for every possible state


(because can't predict where one will end up)

Optimal policy and state values for the given $R(i)$:

3	→	→	→	<div>+ 1</div>
2	↑		↑	<div>- 1</div>
1	↑	←	←	←
	1	2	3	4


3	0.812	0.868	0.912	<div>+ 1</div>
2	0.762		0.660	<div>- 1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

MDP Problem



Given an environment **model as a MDP** create a **policy** for acting that maximizes **lifetime reward**

MDP Problem: Model



Given an environment model as a MDP create a **policy** for acting that maximizes **lifetime reward**

Markov Decision Processes (MDPs)



Model:

- Finite set of states, S
- Finite set of actions, A
- (Probabilistic) state transitions, $\delta(s,a)$
- Reward for each state and action, $R(s,a)$

Process:

- Observe state s_t in S
- Choose action a_t in A
- Receive immediate reward r_t
- State changes to s_{t+1}

Example:




- Legal transitions shown
- Reward on unlabeled transitions is 0.

MDP Environment Assumptions


- Markov Assumption:
Next state and reward is a function only of the current state and action:
 - $s_{t+1} = \delta(s_t, a_t)$
 - $r_t = r(s_t, a_t)$
- Uncertain and Unknown Environment:
 δ and r may be nondeterministic and unknown

MDP Problem: Model



Given an environment model as a MDP create a **policy** for acting that maximizes **lifetime reward**

MDP Problem: Lifetime Reward



Given an environment model as a **MDP** create a **policy** for acting that maximizes lifetime reward

Utility (aka value)

In *sequential* decision problems, preferences are expressed between *sequences* of states

Usually use an *additive* utility function:

$$U([s_1, s_2, s_3, \dots, s_n]) = R(s_1) + R(s_2) + R(s_3) + \dots + R(s_n)$$

(cf. path cost in search problems)

Utility of a *state* (a.k.a. its *value*) is defined to be



$$U(s_i) = \frac{\text{expected sum of rewards until termination}}{\text{assuming optimal actions}}$$

Given the utilities of the states, choosing the best action is just MEU: choose the action such that the expected utility of the immediate successors is highest.

Lifetime Reward

- Finite horizon:
 - Rewards accumulate for a fixed period.
 - $\$100\text{K} + \$100\text{K} + \$100\text{K} = \300K
- Infinite horizon:
 - Assume reward accumulates for ever
 - $\$100\text{K} + \$100\text{K} + \dots = \text{infinity}$
- Discounting:
 - Future rewards not worth as much (a bird in hand ...)
 - Introduce discount factor γ
 $\$100\text{K} + \gamma \$100\text{K} + \gamma^2 \$100\text{K} \dots$ converges
 - Will make the math work



MDP Problem: Lifetime Reward



Given an environment **model as a MDP** create a **policy** for acting that maximizes **lifetime reward**

$$V = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$$

MDP Problem: Policy



Given an environment model as a **MDP** create a policy for acting that maximizes **lifetime reward**

$$V = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$$


Assume deterministic world

Policy $\pi: S \rightarrow A$

- Selects an action for each state.




Optimal policy $\pi^*: S \rightarrow A$

- Selects action for each state that maximizes lifetime reward.



Note: with infinite horizon, policy is stationary and independent of start state

- There are many policies, not all are necessarily optimal.
- There may be several optimal policies.



A sequential decision problem for a fully Observable stochastic environment with Markovian transition model and additive Rewards is called an MDP

Markov Decision Processes

- Motivation
- Markov Decision Processes
- Computing Policies From a Model
 - Value Functions
 - Mapping Value Functions to Policies
 - Computing Value Functions through Value Iteration
 - An Alternative: Policy Iteration (appendix)
- Summary

Value Function V^π for a Given Policy π


- $V^\pi(s_t)$ is the accumulated lifetime reward resulting from starting in state s_t and repeatedly executing policy π :

$$V^\pi(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} \dots$$

$$V^\pi(s_t) = \sum_i \gamma^i r_{t+i}$$

where $r_t, r_{t+1}, r_{t+2} \dots$ are generated by following π , starting at s_t .

Assume $\gamma = .9$



An Optimal Policy π^* Given Value Function V^*

Idea: Given state s

1. Examine **all** possible actions a_i in state s .
2. Select action a_i with greatest lifetime reward.


Lifetime reward $Q(s, a_i)$ is:

- the immediate reward for taking action $r(s, a)$...
- plus lifetime reward starting in target state $V(\delta(s, a))$...
- discounted by γ .

$$\pi^*(s) = \operatorname{argmax}_a [r(s, a) + \gamma V^*(\delta(s, a))]$$

Must Know:

- Value function
- Environment model.
 - $\delta : S \times A \rightarrow S$
 - $r : S \times A \rightarrow \mathbb{R}$




Example: Mapping Value Function to Policy

- Agent selects optimal action from V :

$$\pi(s) = \operatorname{argmax}_a [r(s,a) + \gamma V(\delta(s, a))]$$

Model + V :

$\gamma = 0.9$




Example: Mapping Value Function to Policy

- Agent selects optimal action from V :

$$\pi(s) = \operatorname{argmax}_a [r(s,a) + \gamma V(\delta(s, a))]$$


Model + V :



$\gamma = 0.9$

- $a: 0 + 0.9 \times 100 = 90$
- $b: 0 + 0.9 \times 81 = 72.9$
- select a

π :




Example: Mapping Value Function to Policy

- Agent selects optimal action from V :

$$\pi(s) = \operatorname{argmax}_a [r(s,a) + \gamma V(\delta(s, a))]$$


Model + V :



$\gamma = 0.9$

- $a: 100 + 0.9 \times 0 = 100$
 - $b: 0 + 0.9 \times 90 = 81$
- select a

π :




Example: Mapping Value Function to Policy

- Agent selects optimal action from V :

$$\pi(s) = \operatorname{argmax}_a [r(s,a) + \gamma V(\delta(s, a))]$$


Model + V :



$\gamma = 0.9$

- a: ?
- b: ?
- c: ?
- select ?

π :




Markov Decision Processes

- Motivation
- Markov Decision Processes
- Computing Policies From a Model
 - Value Functions
 - Mapping Value Functions to Policies
 - Computing Value Functions through Value Iteration
 - An Alternative: Policy Iteration
- Summary


Value Function V^* for an optimal policy π^*

Example




- Optimal value function for a one step horizon:

$$V^*_1(s) = \max_{a_i} [r(s, a_i)]$$



Value Function V^* for an optimal policy π^*

Example



- Optimal value function for a one step horizon:


$$V^*_1(s) = \max_{a_i} [r(s, a_i)]$$

- Optimal value function for a two step horizon:

$$V^*_2(s) = \max_{a_i} [r(s, a_i) + \gamma V^*_1(\delta(s, a_i))]$$


Instance of the Dynamic Programming Principle:

- Reuse shared sub-results
- Exponential saving



Value Function V^* for an optimal policy π^*

Example



- Optimal value function for a one step horizon:

$$V^*_1(s) = \max_{a_i} [r(s, a_i)]$$

- Optimal value function for a two step horizon:


$$V^*_2(s) = \max_{a_i} [r(s, a_i) + \gamma V^*_1(\delta(s, a_i))]$$

- Optimal value function for an n step horizon:

$$V^*_n(s) = \max_{a_i} [r(s, a_i) + \gamma V^*_{n-1}(\delta(s, a_i))]$$

Value Function V^* for an optimal policy π^*

Example



- Optimal value function for a one step horizon:

$$V^*_1(s) = \max_{a_i} [r(s, a_i)]$$

- Optimal value function for a two step horizon:

$$V^*_2(s) = \max_{a_i} [r(s, a_i) + \gamma V^*_1(\delta(s, a_i))]$$

- Optimal value function for an n step horizon:

$$V^*_n(s) = \max_{a_i} [r(s, a_i) + \gamma V^*_{n-1}(\delta(s, a_i))]$$

- Optimal value function for an infinite horizon:

$$V^*(s) = \max_{a_i} [r(s, a_i) + \gamma V^*(\delta(s, a_i))]$$

Bellman equation

Definition of utility of states leads to a simple relationship among utilities of neighboring states:

$$\begin{aligned} & \text{expected sum of rewards} \\ &= \text{current reward} \\ &+ \text{expected sum of rewards after taking best action} \end{aligned}$$

Bellman equation (1957):

Model $M_{ij}^a \equiv P(j|i, a)$ = probability that doing a in i leads to j

$$U(i) = R(i) + \max_a \sum_j U(j) M_{ij}^a$$

$$U(1, 1) = -0.04$$

$$\begin{aligned} &+ \max \{ 0.8U(1, 2) + 0.1U(2, 1) + 0.1U(1, 1), && \text{up} \\ &0.9U(1, 1) + 0.1U(1, 2) && \text{left} \\ &0.9U(1, 1) + 0.1U(2, 1) && \text{down} \\ &0.8U(2, 1) + 0.1U(1, 2) + 0.1U(1, 1) \} && \text{right} \end{aligned}$$

One equation per state = n nonlinear equations in n unknowns

Value iteration algorithm


Idea: Start with arbitrary utility values

Update to make them locally consistent with Bellman eqn.

Everywhere locally consistent \Rightarrow global optimality

repeat until “no change”

$$U(i) \leftarrow R(i) + \max_a \sum_j U(j) M_{ij}^a \quad \text{for all } i$$



Solving MDPs by Value Iteration

Insight: Can calculate optimal values iteratively using Dynamic Programming.

Algorithm:

- Iteratively calculate value using Bellman's Equation:

$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

- Terminate when values are “close enough”

$$|V_{t+1}^*(s) - V_t^*(s)| < \varepsilon$$

- Agent selects optimal action by one step lookahead on V^* :

$$\pi^*(s) = \operatorname{argmax}_a [r(s,a) + \gamma V^*(\delta(s, a))]$$

Convergence of Value Iteration

- If terminate when values are “close enough”

$$|V_{t+1}(s) - V_t(s)| < \varepsilon$$

Then:

$$\text{Max}_{s \text{ in } S} |V_{t+1}(s) - V^*(s)| < 2\varepsilon\gamma/(1 - \gamma)$$


- Converges in polynomial time.
- Convergence guaranteed even if updates are performed infinitely often, but asynchronously and in any order.

Example of Value Iteration


$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

$$\gamma = 0.9$$

V_t^*



V_{t+1}^*




- a: $0 + 0.9 \times 0 = 0$
- b: $0 + 0.9 \times 0 = 0$
- Max = 0

Example of Value Iteration


$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

$$\gamma = 0.9$$

V_t^*



V_{t+1}^*




- a: $100 + 0.9 \times 0 = 100$
- b: $0 + 0.9 \times 0 = 0$
- c: $0 + 0.9 \times 0 = 0$
- Max = 100

Example of Value Iteration


$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

$$\gamma = 0.9$$

V_t^*



V_{t+1}^*

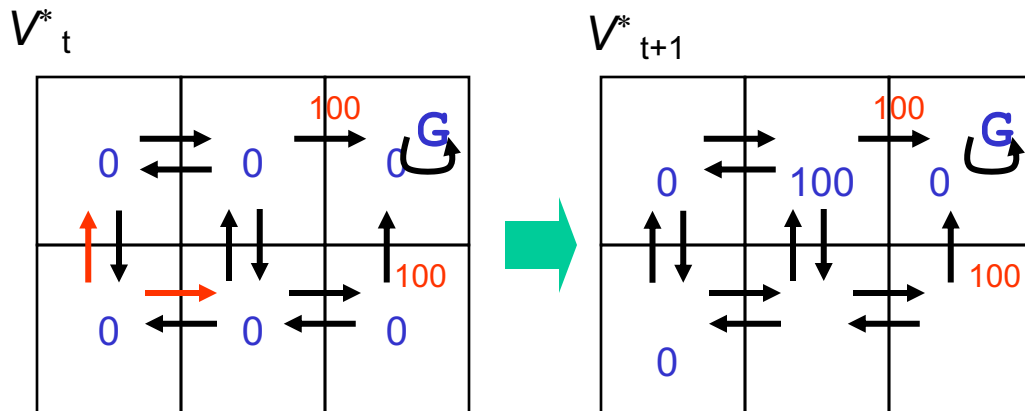


- a: $0 + 0.9 \times 0 = 0$
- Max = 0

Example of Value Iteration

$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

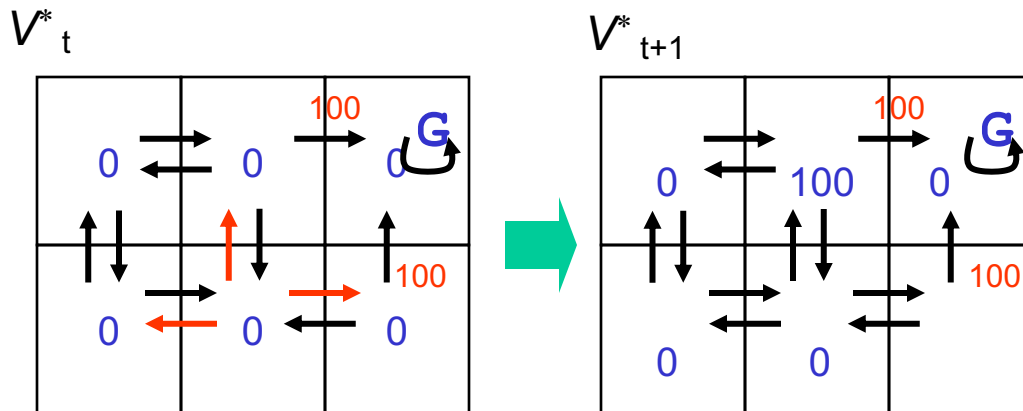
$$\gamma = 0.9$$



Example of Value Iteration

$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

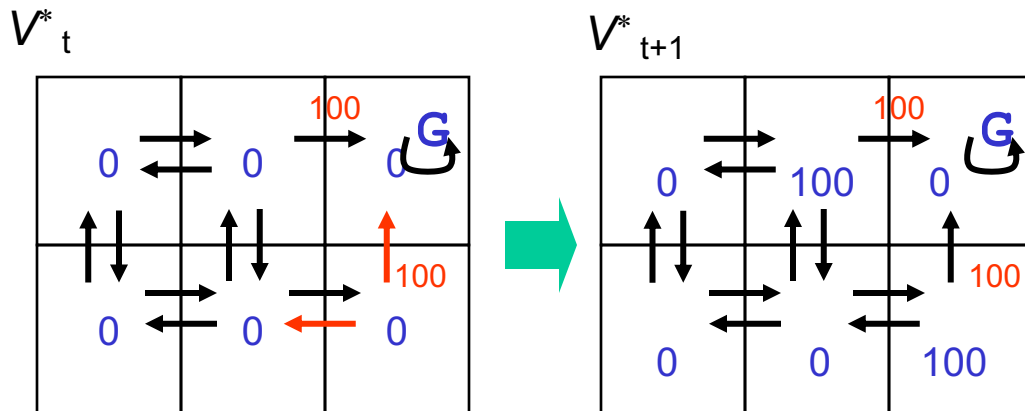
$$\gamma = 0.9$$



Example of Value Iteration

$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

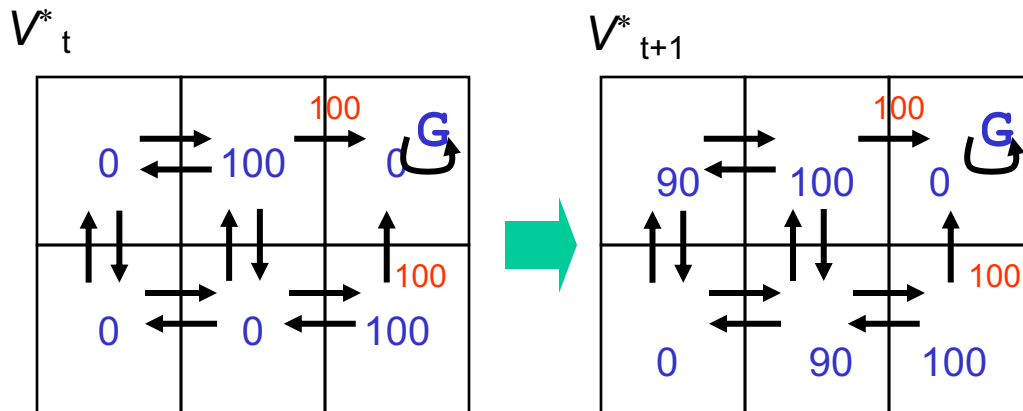
$$\gamma = 0.9$$



Example of Value Iteration

$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

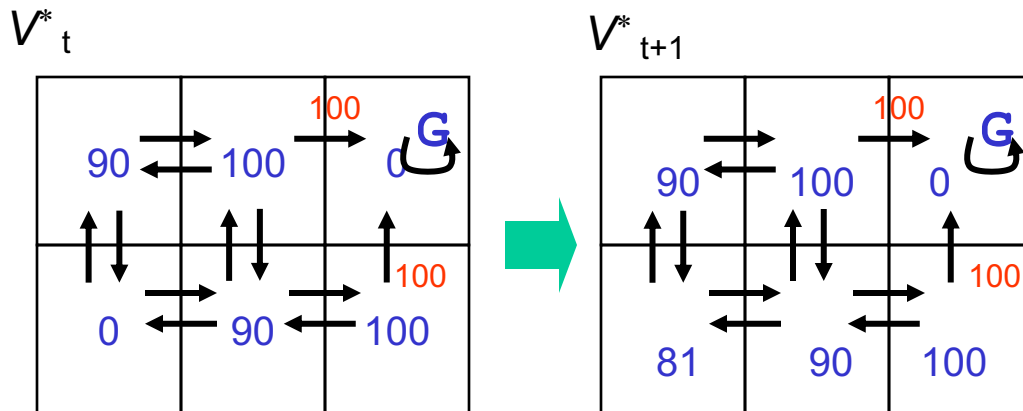
$$\gamma = 0.9$$



Example of Value Iteration

$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

$$\gamma = 0.9$$




Example of Value Iteration


$$V_{t+1}^*(s) \leftarrow \max_a [r(s,a) + \gamma V_t^*(\delta(s, a))]$$

$$\gamma = 0.9$$

V_t^*



V_{t+1}^*



Markov Decision Processes

- Motivation
- Markov Decision Processes
- Computing policies from a modelValue Functions
 - Mapping Value Functions to Policies
 - Computing Value Functions through Value Iteration
 - An Alternative: Policy Iteration
- Summary

Policy iteration

Idea: search for optimal policy and utility values simultaneously

Algorithm:

$\pi \leftarrow$ an arbitrary initial policy

repeat until no change in π

 compute utilities given π

 update π as if utilities were correct (i.e., local MEU)

To compute utilities given a fixed π :

$$U(i) = R(i) + \sum_j U(j) M_{ij}^{\pi(i)} \quad \text{for all } i$$

i.e., n simultaneous linear equations in n unknowns, solve in $O(n^3)$

- Why use policy iteration? May converge faster; convergence guarantees.

Policy Iteration

Idea: Iteratively improve the policy

1. Policy Evaluation: Given a policy π_i calculate $V_i = V^{\pi_i}$, the utility of each state if π_i were to be executed.
2. Policy Improvement: Calculate a new maximum expected utility policy π_{i+1} using one-step look ahead based on V_i .

- π_i improves at every step, converging if $\pi_i \neq \pi_{i+1}$.
- Computing V_i is simpler than for Value iteration (no max):

$$V_{t+1}^*(s) \leftarrow r(s, \pi_i(s)) + \gamma V_t^*(\delta(s, \pi_i(s)))$$

- Solve linear equations in $O(N^3)$
- Solve iteratively, similar to value iteration.