

CSCI-567: Machine Learning (Fall 2019)

Prof. Victor Adamchik

U of Southern California

Nov. 14, 2019

November 14, 2019 1 / 63

General EM algorithm

EM is an algorithm to solve MLE with latent variables (not just GMM), i.e. find the maximizer of

$$P(\theta) = \sum_{n=1}^N \ln p(\mathbf{x}_n; \theta)$$

Directly solving the objective is intractable. Instead we optimize the lower bound

$$P(\theta) \geq F(\theta, \{q_n^{(t)}\})$$

where

$$F(\theta, \{q_n^{(t)}\}) = \sum_{n=1}^N \sum_{k=1}^K \left(q_n(k) \ln p(\mathbf{x}_n, z_n = k; \theta^{(t)}) - q_n(k) \ln q_n(k) \right)$$

November 14, 2019 3 / 63

Outline

- 1 Review of the last lecture
- 2 Density estimation
- 3 Naive Bayes Revisited
- 4 Markov chain
- 5 Hidden Markov Model

November 14, 2019 2 / 63

General EM algorithm

Step 0 Initialize $\theta^{(1)}$, $t = 1$

Step 1 (E-Step) update the posterior of latent variables

$$q_n^{(t)}(z_n = k) = p(z_n = k | \mathbf{x}_n; \theta^{(t)})$$

and obtain **Expectation** of complete likelihood

$$Q(\theta; \theta^{(t)}) = \sum_{n=1}^N \mathbb{E}_{z_n \sim q_n^{(t)}} [\ln p(\mathbf{x}_n, z_n; \theta)]$$


Step 2 (M-Step) update the model parameter via **Maximization**

$$\theta^{(t+1)} \leftarrow \arg\max_{\theta} Q(\theta; \theta^{(t)})$$

Step 3 $t \leftarrow t + 1$ and return to Step 1 if not converged

November 14, 2019 4 / 63

Pictorial explanation



$P(\theta)$ is non-concave, but $Q(\theta; \theta^{(t)})$ often is concave and easy to maximize.

$$\begin{aligned} P(\theta^{(t+1)}) &\geq F(\theta^{(t+1)}; \{q_n^{(t)}\}) \\ &\geq F(\theta^{(t)}; \{q_n^{(t)}\}) \\ &= P(\theta^{(t)}) \end{aligned}$$

So EM always increases the objective value and will converge to some local maximum (similar to K-means).

November 14, 2019 5 / 63

Apply EM to learn GMMs

E-Step:

$$q_n^{(t)}(z_n = k) = p(z_n = k | \mathbf{x}_n; \theta^{(t)})$$

This computes the “soft assignment” $\gamma_{nk} = q_n^{(t)}(z_n = k)$, i.e. conditional probability of \mathbf{x}_n belonging to cluster k .

November 14, 2019 6 / 63

Apply EM to learn GMMs

M-Step:

$$\begin{aligned} \operatorname{argmax}_{\theta} Q(\theta, \theta^{(t)}) &= \operatorname{argmax}_{\theta} \sum_{n=1}^N \mathbb{E}_{z_n \sim q_n^{(t)}} [\ln p(\mathbf{x}_n, z_n; \theta)] \\ &= \operatorname{argmax}_{\theta} \sum_{n=1}^N \mathbb{E}_{z_n \sim q_n^{(t)}} [\ln p(z_n; \theta) + \ln p(\mathbf{x}_n | z_n; \theta)] \\ &= \operatorname{argmax}_{\{\omega_k, \mu_k, \Sigma_k\}} \sum_{n=1}^N \sum_{k=1}^K \gamma_{nk} (\ln \omega_k + \ln N(\mathbf{x}_n | \mu_k, \Sigma_k)) \end{aligned}$$

To find $\omega_1, \dots, \omega_K$, solve

$$\operatorname{argmax}_{\omega} \sum_{n=1}^N \sum_{k=1}^K \gamma_{nk} \ln \omega_k$$

To find each μ_k, Σ_k , solve

$$\operatorname{argmax}_{\mu_k, \Sigma_k} \sum_{n=1}^N \gamma_{nk} \ln N(\mathbf{x}_n | \mu_k, \Sigma_k)$$

November 14, 2019 7 / 63

M-Step (continued)

Solutions to previous two problems are very natural, for each k

$$\omega_k = \frac{\sum_n \gamma_{nk}}{N}$$

i.e. (weighted) fraction of examples belonging to cluster k

$$\mu_k = \frac{\sum_n \gamma_{nk} \mathbf{x}_n}{\sum_n \gamma_{nk}}$$

i.e. (weighted) average of examples belonging to cluster k

$$\Sigma_k = \frac{1}{\sum_n \gamma_{nk}} \sum_n \gamma_{nk} (\mathbf{x}_n - \mu_k)(\mathbf{x}_n - \mu_k)^T$$

i.e (weighted) covariance of examples belonging to cluster k

November 14, 2019 8 / 63

GMM: putting it together

EM for clustering:

Step 0 Initialize $\omega_k, \mu_k, \Sigma_k$ for each $k \in [K]$

Step 1 (E-Step) update the “soft assignment” (fixing parameters)

$$\gamma_{nk} = p(z_n = k \mid \mathbf{x}_n)$$

Step 2 (M-Step) update the model parameter (fixing assignments)

$$\omega_k = \frac{\sum_n \gamma_{nk}}{N} \quad \mu_k = \frac{\sum_n \gamma_{nk} \mathbf{x}_n}{\sum_n \gamma_{nk}}$$
$$\Sigma_k = \frac{1}{\sum_n \gamma_{nk}} \sum_n \gamma_{nk} (\mathbf{x}_n - \mu_k)(\mathbf{x}_n - \mu_k)^T$$

Step 3 return to Step 1 if not converged

November 14, 2019 9 / 63

Outline

- 1 Review of the last lecture
- 2 **Density estimation**
- 3 Naive Bayes Revisited
- 4 Markov chain
- 5 Hidden Markov Model

November 14, 2019 11 / 63

Connection to K-means

K-means is in fact a **special case** of EM for (a simplified) GMM:

Let $\Sigma_k = \sigma^2 \mathbf{I}$ for some fixed σ , so only ω_k and μ_k are parameters.

EM becomes K-means:

$$\operatorname{argmax}_{\theta} \prod_{n=1}^N p(\mathbf{x}_n; \theta) = \operatorname{argmax}_{\theta} \prod_{n=1}^N \sum_{k=1}^K p(z_n = k) N(\mathbf{x}_n | \mu_k)$$

If we assume hard assignments $p(z_n = k) = 1$, if $k = C(n)$, then

$$\operatorname{argmax}_{\theta} \prod_{n=1}^N p(\mathbf{x}_n; \theta) = \operatorname{argmax}_{\theta} \prod_{n=1}^N N(\mathbf{x}_n | \mu_{C(n)})$$
$$= \operatorname{argmax}_{\theta} \prod_{n=1}^N \exp\left(\frac{-1}{2\sigma^2} \|\mathbf{x}_n - \mu_{C(n)}\|_2^2\right) = \operatorname{argmin}_{\mu, C} \sum_{n=1}^N \|\mathbf{x}_n - \mu_{C(n)}\|_2^2$$

GMM is a soft version of K-means and it provides a probabilistic interpretation of the data.

November 14, 2019 10 / 63

Density estimation

Observe what we have done indirectly for clustering with GMMs is:

Given a training set $\mathbf{x}_1, \dots, \mathbf{x}_N$, **estimate a density function p that could have generated this dataset** (via $\mathbf{x}_n \stackrel{i.i.d.}{\sim} p$).

We say that a random variable x has a probability distribution $p(x)$.

This is exactly the problem of **density estimation**, another important unsupervised learning problem.

Useful for many downstream applications

- we have seen clustering already, will see more applications today
- these applications also **provide a way to measure quality of the density estimator**

November 14, 2019 12 / 63

Parametric generative models

Parametric estimation assumes a generative model parametrized by θ :

$$p(\mathbf{x}) = p(\mathbf{x} ; \theta)$$

here $p(x)$ is a common (predefined) probability distribution. Examples:

- **GMM**: $p(\mathbf{x} ; \theta) = \sum_{k=1}^K \omega_k N(\mathbf{x} \mid \mu_k, \Sigma_k)$ where $\theta = \{\omega_k, \mu_k, \Sigma_k\}$
- **Multinomial** for 1D examples with K possible values

$$p(x = k ; \theta) = \theta_k$$

where θ is a distribution over K elements.

Size of θ is independent of the training set size, so it's **parametric**.

MLE for multinomial

$$\begin{aligned} \operatorname{argmax}_{\theta} &= \sum_{n=1}^N \ln p(x = x_n ; \theta) = \sum_{n=1}^N \ln \theta_{x_n} \\ &= \sum_{k=1}^K \sum_{n: x_n=k} \ln \theta_k = \sum_{k=1}^K z_k \ln \theta_k \end{aligned}$$

where $z_k = |\{n : x_n = k\}|$ is the number of examples with value k .

The solution (verify yourself!) is simply

$$\theta_k = \frac{z_k}{N} \propto z_k,$$

i.e. the fraction of examples with value k .

Parametric methods

Again, we apply **MLE** to learn the parameters θ :

$$\operatorname{argmax}_{\theta} = \sum_{n=1}^N \ln p(x_n ; \theta)$$

For some cases this is intractable and we can use **EM** to approximately solve MLE (e.g. GMMs).

For some other cases this admits a simple closed-form solution (e.g. multinomial).

Nonparametric models

Can we estimate **without** assuming a fixed generative model?

Kernel density estimation (KDE) is a common approach for nonparametric density estimation (without a pre-defined distribution).

Here “kernel” means something different from what we have seen for “kernel function”.

The scikit-learn library provides the `KernelDensity` class that implements KDE.


We focus on the 1D (continuous) case.

High level idea

picture from Wikipedia

KDE is closely related to a **histogram**. A histogram is a plot that involves first grouping the observations into bins and counting the number of events that fall into each bin. To construct KDE,

- for each data point, create a “hump” (via a kernel)
- sum up all the humps; more data - a higher hump



November 14, 2019 17 / 63

Kernel


KDE with a kernel $K(x): \mathbb{R} \rightarrow \mathbb{R}$ centered at x_n :

$$p(x) = \frac{1}{N} \sum_{n=1}^N K(x - x_n)$$

Many choices for K , for example, $K(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}}$, the **standard Gaussian density**

Properties of a kernel:

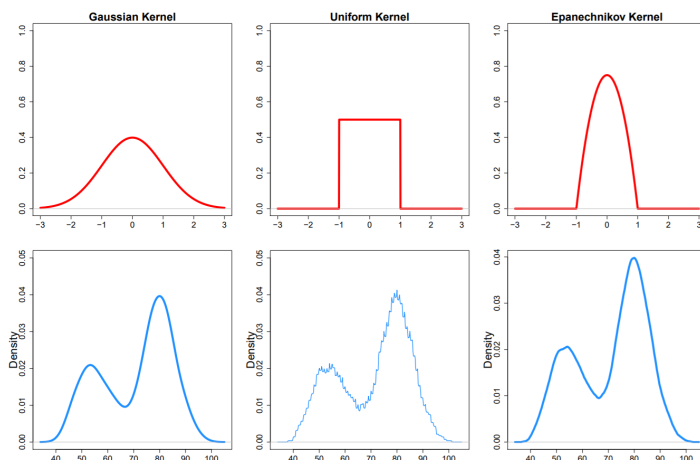
- **symmetry**: $K(x) = K(-x)$
- $\int_{-\infty}^{\infty} K(x) dx = 1$, this insures p is a **density function**.



November 14, 2019 18 / 63

Different kernels $K(x)$

$$\frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}}, \quad \frac{1}{2} \mathbb{I}[|x| \leq 1], \quad \frac{3}{4} \max\{1 - x^2, 0\}$$



November 14, 2019 19 / 63

Bandwidth

If $K(x)$ is a kernel, then for any $h > 0$

$$K_h(u) \triangleq \frac{1}{h} K\left(\frac{u}{h}\right) \quad (\text{stretching the kernel})$$

can be used as a kernel too (verify the two properties yourself)

So, general KDE is determined by both the kernel K and the bandwidth h

$$p(x) = \frac{1}{N} \sum_{n=1}^N K_h(x - x_n) = \frac{1}{Nh} \sum_{n=1}^N K\left(\frac{x - x_n}{h}\right)$$

- x_n controls the **center** of each hump
- h controls the **width/variance** of the humps

November 14, 2019 20 / 63


A larger h will smooth a density.

A small h will yield a density that is spiky and very hard to interpret.

Assume Gaussian kernel.

Gray curve is ground-truth

- Red: $h = 0.05$
- Black: $h = 0.337$
- Green: $h = 2$



Summary

This was a gentle introduction to probability density estimation.

- Histogram provides a fast and reliable way to visualize the probability density of data.
- Parametric probability density estimation involves selecting a common distribution and estimating the parameters for the density function from data.
- Nonparametric probability density estimation involves using an algorithm (KDE) to fit a model to the arbitrary distribution of data.

Selecting h is a deep topic

- one can also do **cross-validation** based on downstream applications
- there are theoretically-motivated approaches

Find a value of h that minimizes the error between the estimated density and the true density:

$$\mathbb{E} [(p_{KDE}(x) - p(x))^2] = \mathbb{E} [p_{KDE}(x) - p(x)]^2 + Var [p_{KDE}(x)]$$

This expression is an example of the bias-variance tradeoff, which we saw in the earlier lecture.

Outline

- 1 Review of the last lecture
- 2 Density estimation
- 3 Naive Bayes Revisited
- 4 Markov chain
- 5 Hidden Markov Model

Bayes optimal classifier

Suppose the data (\mathbf{x}_n, y_n) is drawn from a joint distribution $p(\mathbf{x}, y)$, the **Bayes optimal classifier** is

$$f^*(\mathbf{x}) = \operatorname{argmax}_{c \in [C]} p(c | \mathbf{x})$$

i.e. **predict the class with the largest conditional probability**.

$p(\mathbf{x}, y)$ is of course unknown, but we can estimate it, which is **exactly a density estimation problem!**

Observe that

$$p(\mathbf{x}, y) = p(y)p(\mathbf{x} | y)$$

To estimate $p(\mathbf{x} | y = c)$ for some $c \in [C]$, we are doing density estimation using data with label $y = c$.

November 14, 2019 25 / 63

Continuous features

If the feature is continuous, we can do

- **parametric estimation**, e.g. via a Gaussian

$$p(x_d = x | y = c) = \frac{1}{\sqrt{2\pi}\sigma_{cd}} \exp\left(-\frac{(x - \mu_{cd})^2}{2\sigma_{cd}^2}\right)$$

where μ_{cd} and σ_{cd}^2 are the empirical mean and variance of feature d among all examples with label c .

- or **nonparametric estimation**, e.g. via a kernel K and bandwidth h :

$$p(x_d = x | y = c) = \frac{1}{|\{n : y_n = c\}|} \sum_{n: y_n = c} K_h(x - x_{nd})$$

November 14, 2019 27 / 63

Discrete features

For a label $c \in [C]$,

$$p(y = c) = \frac{|\{n : y_n = c\}|}{N}$$

For each possible value k of a discrete feature d ,

$$p(x_d = k | y = c) = \frac{|\{n : x_{nd} = k, y_n = c\}|}{|\{n : y_n = c\}|}$$

November 14, 2019 26 / 63

How to predict?

Using Naive Bayes assumption:

$$p(\mathbf{x} | y = c) = \prod_{d=1}^D p(x_d | y = c)$$

the **prediction** for a new example \mathbf{x} is

$$\begin{aligned} \operatorname{argmax}_{c \in [C]} p(y = c | \mathbf{x}) &= \operatorname{argmax}_{c \in [C]} \frac{p(\mathbf{x} | y = c)p(y = c)}{p(\mathbf{x})} \\ &= \operatorname{argmax}_{c \in [C]} \left(p(y = c) \prod_{d=1}^D p(x_d | y = c) \right) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln p(y = c) + \sum_{d=1}^D \ln p(x_d | y = c) \right) \end{aligned}$$

November 14, 2019 28 / 63

Naive Bayes

For **discrete features**, plugging in previous MLE estimations gives

$$\begin{aligned} & \operatorname{argmax}_{c \in [C]} p(y = c \mid \mathbf{x}) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln p(y = c) + \sum_{d=1}^D \ln p(x_d \mid y = c) \right) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln |\{n : y_n = c\}| + \sum_{d=1}^D \ln \frac{|\{n : x_{nd} = x_d, y_n = c\}|}{|\{n : y_n = c\}|} \right) \end{aligned}$$

Connection to logistic regression

Let us fix the variance for each feature to be σ (i.e. not a parameter of the model any more), then the prediction becomes

$$\begin{aligned} & \operatorname{argmax}_{c \in [C]} p(y = c \mid \mathbf{x}) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln |\{n : y_n = c\}| - \sum_{d=1}^D \left(\ln \sigma + \frac{(x_d - \mu_{cd})^2}{2\sigma^2} \right) \right) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln |\{n : y_n = c\}| - \frac{\|\mathbf{x}\|_2^2}{2\sigma^2} - \sum_{d=1}^D \frac{\mu_{cd}^2}{2\sigma^2} + \sum_{d=1}^D \frac{\mu_{cd}}{\sigma^2} x_d \right) \\ &= \operatorname{argmax}_{c \in [C]} \left(w_{c0} + \sum_{d=1}^D w_{cd} x_d \right) = \operatorname{argmax}_{c \in [C]} \mathbf{w}_c^T \mathbf{x} \quad (\text{linear classifier!}) \end{aligned}$$

where we denote $w_{c0} = \ln |\{n : y_n = c\}| - \sum_{d=1}^D \frac{\mu_{cd}^2}{2\sigma^2}$ and $w_{cd} = \frac{\mu_{cd}}{\sigma^2}$.

Naive Bayes

For **continuous features** with a Gaussian model,

$$\begin{aligned} & \operatorname{argmax}_{c \in [C]} p(y = c \mid \mathbf{x}) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln p(y = c) + \sum_{d=1}^D \ln p(x_d \mid y = c) \right) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln |\{n : y_n = c\}| + \sum_{d=1}^D \ln \left(\frac{1}{\sqrt{2\pi}\sigma_{cd}} \exp \left(-\frac{(x_d - \mu_{cd})^2}{2\sigma_{cd}^2} \right) \right) \right) \\ &= \operatorname{argmax}_{c \in [C]} \left(\ln |\{n : y_n = c\}| - \sum_{d=1}^D \left(\ln \sigma_{cd} + \frac{(x_d - \mu_{cd})^2}{2\sigma_{cd}^2} \right) \right) \end{aligned}$$

Connection to logistic regression

You can verify

$$p(y = c \mid \mathbf{x}) \propto e^{\mathbf{w}_c^T \mathbf{x}}$$

This is exactly the **softmax** function, the same model we used for a probabilistic interpretation of logistic regression!

So what is different then? They **learn the parameters in different ways**:

- both via MLE, **one** on $p(y = c \mid \mathbf{x})$, **the other** on $p(\mathbf{x}, y)$
- solutions are different: **logistic regression has no closed-form**, **naive Bayes admits a simple closed-form**

Two different modeling paradigms

Suppose the training data is from an *unknown* joint probabilistic model $p(\mathbf{x}, y)$. There are two kinds of classification models in machine learning — *generative* models and *discriminative* models.

Differences in *assuming* models for the data

- the generative approach requires we specify the model for the joint distribution (such as Naive Bayes), and thus, maximize the *joint* likelihood $\sum_n \log p(\mathbf{x}_n, y_n)$
- the discriminative approach (discriminative) requires only specifying a model for the conditional distribution (such as logistic regression), and thus, maximize the *conditional* likelihood $\sum_n \log p(y_n | \mathbf{x}_n)$
- Sometimes, modeling by discriminative approach is easier
- Sometimes, parameter estimation by generative approach is easier


November 14, 2019 33 / 63

Generative model v.s discriminative model

	Discriminative model	Generative model
Example	logistic regression	naive Bayes
Model	conditional $p(y x)$	joint $p(x, y)$ (might have same $p(y x)$)
Learning	MLE	MLE
Accuracy	usually better for large N	usually better for small N
Remark		more flexible, can generate data after learning

November 14, 2019 34 / 63

Determining sex (man or woman) based on measurements




November 14, 2019 35 / 63

Example: Generative approach

Propose a model of the joint distribution of ($x = \text{height}$, $y = \text{sex}$)

our data

Sex	Height
1	6'
2	5'2"
1	5'6"
1	6'2"
2	5.7"
...	...



Intuition: we will model how heights vary (according to a Gaussian) in each sub-population (male and female).


Note: This is similar to Naive Bayes for detecting spam emails.

November 14, 2019 36 / 63

Model of the joint distribution

$$p(x, y) = p(y)p(x|y)$$
$$= \begin{cases} p_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} & \text{if } y = 1 \\ p_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x-\mu_2)^2}{2\sigma_2^2}} & \text{if } y = 2 \end{cases}$$

where $p_1 + p_2 = 1$ represents two **prior** probabilities that x is given the label 1 or 2 respectively. $p(x|y)$ is assumed to be Gaussians.



Parameter estimation

Likelihood of the training data $\mathcal{D} = \{(x_n, y_n)\}_{n=1}^N$ with $y_n \in \{1, 2\}$

$$\begin{aligned} \log P(\mathcal{D}) &= \sum_n \log p(x_n, y_n) \\ &= \sum_{n: y_n=1} \log \left(p_1 \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x_n-\mu_1)^2}{2\sigma_1^2}} \right) \\ &\quad + \sum_{n: y_n=2} \log \left(p_2 \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(x_n-\mu_2)^2}{2\sigma_2^2}} \right) \end{aligned}$$

Maximize the likelihood function

$$(p_1^*, p_2^*, \mu_1^*, \mu_2^*, \sigma_1^*, \sigma_2^*) = \operatorname{argmax} \log P(\mathcal{D})$$

Decision boundary

The decision boundary between two classes is defined by

$$p(y = 1|x) \geq p(y = 2|x)$$

which is equivalent to

$$p(x|y = 1)p(y = 1) \geq p(x|y = 2)p(y = 2)$$

Namely,


$$-\frac{(x - \mu_1)^2}{2\sigma_1^2} - \log \sqrt{2\pi}\sigma_1 + \log p_1 \geq -\frac{(x - \mu_2)^2}{2\sigma_2^2} - \log \sqrt{2\pi}\sigma_2 + \log p_2$$

It is quadratic in x . It follows (for some a , b and c , that

$$ax^2 + bx + c \geq 0$$

The decision boundary is **not linear**!

Example of nonlinear decision boundary



Note: the boundary is characterized by a quadratic function, giving rise to the shape of parabolic curve.

A special case

What if we assume the two Gaussians have the same variance?

We will get a *linear* decision boundary

From the previous slide:

$$-\frac{(x - \mu_1)^2}{2\sigma_1^2} - \log \sqrt{2\pi}\sigma_1 + \log p_1 \geq -\frac{(x - \mu_2)^2}{2\sigma_2^2} - \log \sqrt{2\pi}\sigma_2 + \log p_2$$

Setting $\sigma_1 = \sigma_2$, we obtain

$$bx + c \geq 0$$

Note: equal variances across two different categories could be a very strong assumption.

For example, the plot suggests that the *male* population has slightly bigger variance (i.e., bigger eclipse) than the *female* population.

November 14, 2019 41 / 63

Markov Models

Markov models are powerful **probabilistic models** to analyze sequential data. A.A.Markov (1856-1922) introduced the Markov chains in 1906 when he produced the first theoretical results for stochastic processes. They are now commonly used in

- text or speech recognition
- stock market prediction
- bioinformatics
- ...

November 14, 2019 43 / 63

Outline

- 1 Review of the last lecture
- 2 Density estimation
- 3 Naive Bayes Revisited
- 4 **Markov chain**
- 5 Hidden Markov Model

November 14, 2019 42 / 63

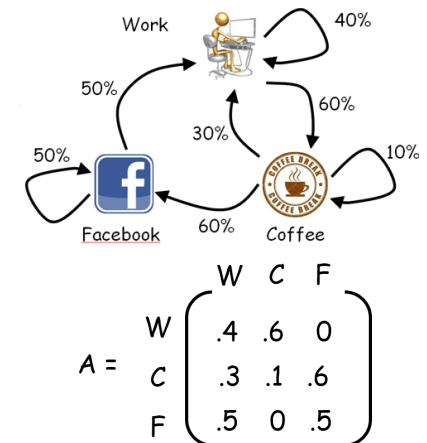
Markov chain

Directed strongly connected graph with self-loops.

Each edge labeled by a positive probability.

At each state, the probabilities on outgoing edges sum up to 1.

Transition (or stochastic) matrix:
 $A = a_{ij} = P(i \rightarrow j \text{ in 1 step}).$



November 14, 2019 44 / 63

Markov chain

Definition

Given a sequentially ordered random variables $X_1, X_2, \dots, X_t, \dots, X_T$, called **states**,

- **Transition probability** for describing how the state at time $t - 1$ changes to the state at time t ,

$$P(X_t = \text{value}' | X_{t-1} = \text{value})$$

- **Initial probability** for describing the initial state at time $t = 1$.

$$P(X_1 = \text{value})$$

All X_t 's take value from the same **discrete** set $\{1, \dots, N\}$.

We will assume that the transition probability does not change with respect to time t , i.e., a stationary Markov chain.

Examples

- Example 1 (**Language model**)

States $[N]$ represent a dictionary of words,

$$a_{\text{ice,cream}} = P(X_{t+1} = \text{cream} | X_t = \text{ice})$$

is an example of the transition probability.

- Example 2 (**Weather**)

States $[N]$ represent weather at each day

$$a_{\text{sunny,rainy}} = P(X_{t+1} = \text{rainy} | X_t = \text{sunny})$$

Markov chain

- Transition probabilities make a table/matrix A whose elements are

$$a_{ij} = P(X_t = j | X_{t-1} = i)$$

- Initial probability becomes a vector π whose elements are

$$\pi_i = P(X_1 = i)$$

where i or j index over from 1 to N . We have the following constraints

$$\sum_j a_{ij} = 1 \quad \sum_i \pi_i = 1$$

Additionally, all those numbers should be non-negative.

Definition

A Markov chain is a stochastic process with the **Markov property**: a sequence of random variables X_1, X_2, \dots s.t.

$$P(X_{t+1} | X_1, X_2, \dots, X_t) = P(X_{t+1} | X_t)$$

i.e. *the current state only depends on the most recent state.*

Is the Markov assumption reasonable? Not completely for the language model for example.

Higher order Markov chains make it more reasonable, e.g.

$$P(X_{t+1} | X_1, X_2, \dots, X_t) = P(X_{t+1} | X_t, X_{t-1})$$

i.e. the current word only depends on the last two words.

Chain Rule

In all derivations we will be using the chain rule:


$$P(X, Y) = P(X | Y) P(Y) = P(Y | X) P(X)$$

$$P(X, Y, Z) = P(X, Y | Z) P(Z)$$

$$P(X, Y, Z) = P(X | Y, Z) P(Y | Z) P(Z)$$

Exercise 1

Consider the following Markov model.
Given that now I am having Coffee, what's the probability that the next step is Facebook and the next is Work?



$$P(X_3 = W, X_2 = F | X_1 = C) =$$

$$\begin{aligned} &= \frac{P(X_3 = W, X_2 = F, X_1 = C)}{P(X_1 = C)} \\ &= \frac{P(X_3 = W | X_2 = F, X_1 = C) P(X_2 = F | X_1 = C) P(X_1 = C)}{P(X_1 = C)} \quad (\text{chain rule}) \\ &= P(X_3 = W | X_2 = F) P(X_2 = F | X_1 = C) \quad (\text{Markov rule}) \\ &= 0.5 \times 0.6 = 0.3 \end{aligned}$$


November 14, 2019 49 / 63

November 14, 2019 50 / 63

Exercise 2

Given that now I am having Coffee, what is the probability that in two steps I am at Work?

$$\begin{aligned} P(X_3 = W | X_1 = C) &= \\ &= \sum_s P(X_3 = W, X_2 = s | X_1 = C) = \end{aligned}$$



$$\begin{aligned} &= P(X_3 = W | X_2 = W) P(X_2 = W | X_1 = C) \quad (\text{marginalization}) \\ &+ P(X_3 = W | X_2 = C) P(X_2 = C | X_1 = C) \\ &+ P(X_3 = W | X_2 = F) P(X_2 = F | X_1 = C) \\ &= 0.3 \times 0.4 + 0.1 \times 0.3 + 0.6 \times 0.5 = 0.45 \end{aligned}$$

Using a transition matrix:

$$P(X_3 = j | X_1 = i) = \sum_{k=1}^N a_{ik} a_{kj} = a_{ij}^2$$

Parameter estimation for Markov models

Now suppose we have observed M sequences of examples:

- $x_{1,1}, \dots, x_{1,T}$
- \dots
- $x_{M,1}, \dots, x_{M,T}$

where

- for simplicity we assume each sequence has the same length T
- lower case $x_{n,t}$ represents the value of the random variable $X_{n,t}$

From these observations how do we *learn the model parameters* (π, \mathbf{A}) ?

November 14, 2019 51 / 63

November 14, 2019 52 / 63

Finding the MLE

Same story, **Maximum Likelihood Estimation**:

$$\operatorname{argmax}_{\pi, \mathbf{A}} \ln P(X_1 = x_1, X_2 = x_2, \dots, X_T = x_T)$$

First, we need to compute this joint probability. Applying the chain rule for random variables, we get

$$\begin{aligned} P(X_1, X_2, \dots, X_T) &= P(X_2, X_3, \dots, X_T | X_1) P(X_1) \\ &= P(X_3, \dots, X_T | X_1, X_2) P(X_2 | X_1) P(X_1) \\ &= \dots = \\ &= P(X_1) \prod_{t=2}^T P(X_t | X_1, \dots, X_{t-1}) \quad (\text{Markov property}) \\ &= P(X_1) \prod_{t=2}^T P(X_t | X_{t-1}) \end{aligned}$$

November 14, 2019 53 / 63

Finding the MLE

The log-likelihood of a sequence x_1, \dots, x_T is

$$\begin{aligned} \ln P(X_1 = x_1, X_2 = x_2, \dots, X_T = x_T) &= \ln P(X_1 = x_1) + \sum_{t=2}^T \ln P(X_t = x_t | X_{t-1} = x_{t-1}) \\ &= \ln \pi_{x_1} + \sum_{t=2}^T \ln a_{x_{t-1}, x_t} \\ &= \sum_n \mathbb{I}[x_1 = n] \ln \pi_n + \sum_{n, n'} \left(\sum_{t=2}^T \mathbb{I}[x_{t-1} = n, x_t = n'] \right) \ln a_{n, n'} \end{aligned}$$

November 14, 2019 54 / 63

Finding the MLE

So MLE is

$$\begin{aligned} \operatorname{argmax}_{\pi, \mathbf{A}} \sum_n (\text{\#initial states with value } n) \ln \pi_n \\ + \sum_{n, n'} (\text{\#transitions from } n \text{ to } n') \ln a_{n, n'} \end{aligned}$$

We have seen this many times. The solution is (derivation is left as an exercise):

$$\begin{aligned} \pi_n &= \frac{\text{\#of sequences starting with } n}{\text{\#of sequences}} \\ a_{n, n'} &= \frac{\text{\#of transitions from } n \text{ to } n'}{\text{\#of transitions starting with } n} \end{aligned}$$


November 14, 2019 55 / 63

Example

Suppose we observed the following 2 sequences of length 5

- sunny, sunny, rainy, rainy, rainy
- rainy, sunny, sunny, sunny, rainy

MLE is the following model



November 14, 2019 56 / 63

Outline

- 1 Review of the last lecture
- 2 Density estimation
- 3 Naive Bayes Revisited
- 4 Markov chain
- 5 Hidden Markov Model

November 14, 2019 57 / 63

Markov Model with outcomes

Now suppose each state X_t also “emits” some **outcome** $O_t \in [O]$ based on the following model

$$P(O_t = o \mid X_t = s) = b_{s,o} \quad (\text{emission probability})$$

independent of anything else.

For example, in the language model, O_t is the speech signal for the underlying word X_t (very useful for **speech recognition**).

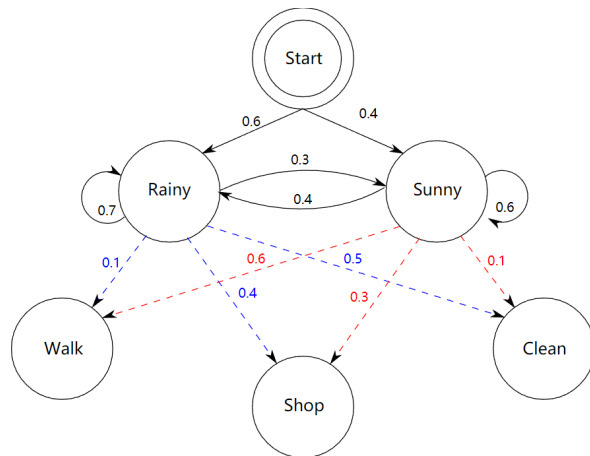
Now the model parameters are $(\{\pi_s\}, \{a_{s,s'}\}, \{b_{s,o}\}) = (\boldsymbol{\pi}, \boldsymbol{A}, \boldsymbol{B})$.

November 14, 2019 58 / 63

Another example

picture from Wikipedia

On each day, we also observe **Bob's activity: walk, shop, or clean**, which only depends on the weather of that day.



November 14, 2019 59 / 63

HMM defines a joint probability

$$\begin{aligned} P(X_1, X_2, \dots, X_T, O_1, O_2, \dots, O_T) \\ = P(X_1, X_2, \dots, X_T) P(O_1, O_2, \dots, O_T \mid X_1, X_2, \dots, X_T) \end{aligned}$$

- Markov assumption simplifies the first term

$$P(X_1, X_2, \dots, X_T) = P(X_1) \prod_{t=2}^T P(X_t \mid X_{t-1})$$

- The *independence* assumption simplifies the second term

$$P(O_1, O_2, \dots, O_T \mid X_1, X_2, \dots, X_T) = \prod_{t=1}^T P(O_t \mid X_t)$$

Namely, each O_t is conditionally independent of anything else, if conditioned on X_t .

November 14, 2019 60 / 63

Joint likelihood

The joint log-likelihood is

$$\begin{aligned} & \ln P(X_1 = x_1, X_2 = x_2, \dots, X_T = x_T, O_1 = o_1, O_2 = o_2, \dots, O_T = o_T) \\ &= \ln P(X_1 = x_1) \prod_{t=2}^T P(X_t = x_t | X_{t-1} = x_{t-1}) \prod_{t=1}^T P(O_t = o_t | X_t = x_t) \\ &= \ln P(X_1 = x_1) + \sum_{t=2}^T \ln P(X_t = x_t | X_{t-1} = x_{t-1}) \\ &\quad + \sum_{t=1}^T \ln P(O_t = o_t | X_t = x_t) \\ &= \ln \pi_{x_1} + \sum_{t=2}^T \ln a_{x_{t-1}, x_t} + \sum_{t=1}^T \ln b_{x_t, o_t} \end{aligned}$$

Learning the model

If we observe M state-outcome sequences: $x_{m,1}, o_{m,1}, \dots, x_{m,T}, o_{m,T}$ for $m = 1, \dots, M$, the MLE is again very simple (verify yourself):

$$\begin{aligned} \pi_s &\propto \text{\#initial states with value } s \\ a_{s,s'} &\propto \text{\#transitions from } s \text{ to } s' \\ b_{s,o} &\propto \text{\#state-outcome pairs } (s, o) \end{aligned}$$

Learning the model

However, *most often we do not observe the states!* Think about the speech recognition example.

This is called **Hidden Markov Model (HMM)**.

Notice that “hidden” is referred to the states of the Markov chain, not to the parameters of the model.

A generic hidden Markov model is illustrated in this picture:

