

# Data Representation & External Sorting

DSCI 551

Wensheng Wu

# Outline

- Representing data



- How are tables stored on storage devices?

- External Sorting

- How to sort 1TB data using 1GB of memory?


# Representing Data Elements

- Relational database elements:

```
CREATE TABLE Product (  
  pid INT PRIMARY KEY,  
  name CHAR(20),  
  description VARCHAR(200),  
  maker CHAR(10) REFERENCES Company(name))
```

- A tuple is stored as a "record"


# Record Formats: Fixed Length



Base address (B)      Address =  $B + L1 + L2$

- Information about field types is the same for all records in a file; stored in *system catalogs*.
- **Note the importance of schema information!**

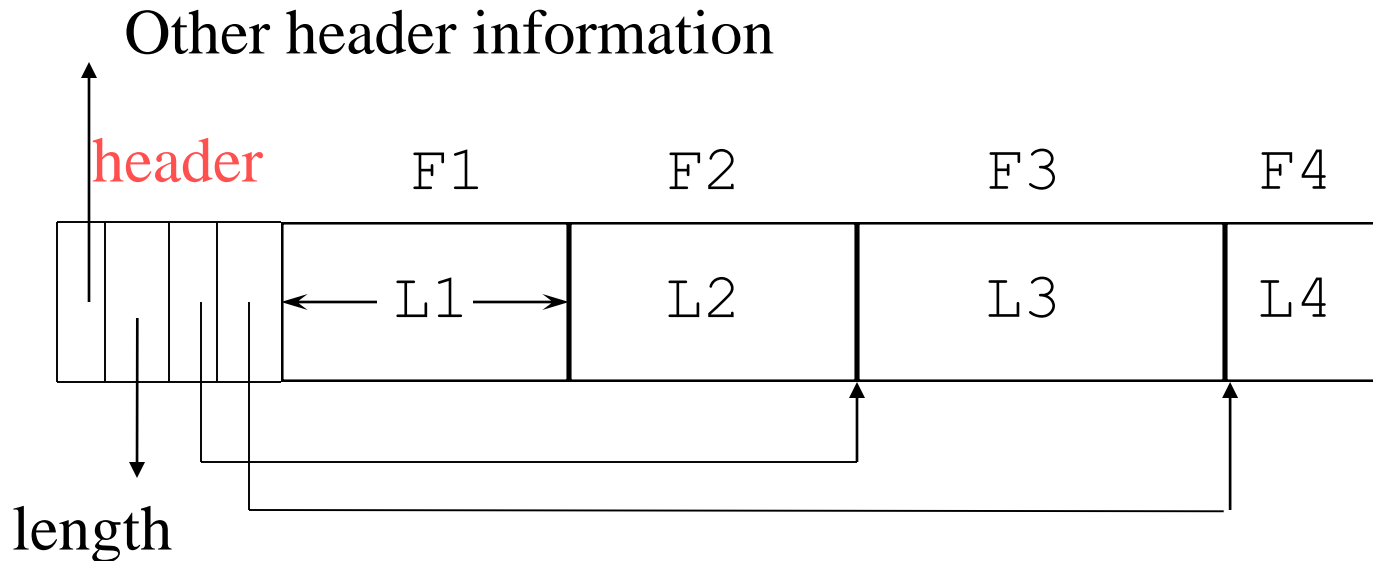
# Record Header



Header:

- Pointer to schema: help finding fields
- Length: so we know where the record ends w/o consulting schema
- Timestamp: time when record last modified or read

# Variable Length Records



Place the fixed fields first: F1, F2


Then the variable length fields: F3, F4

Note: actually no need for pointer to F3, why?

# Storing Records in Blocks

- Blocks have fixed size (typically 4KB)
  - But records may have variable-length

Offset table (slot directory)




Why are records placed from the end?

# Problem with this design?


- Records right after slot directory
- Free space after all records

Offset table (slot directory)





# Outline

- Representing data
  - How are tables stored on storage devices?
- External Sorting 
  - How to sort 1TB data using 1GB of memory?

# The I/O Model of Computation

- In main memory algorithms:
  - we care about CPU time
- In databases
  - time is dominated by I/O cost
- Assumption: cost is given only by I/O
- Consequence: need to redesign certain algorithms, e.g. sorting

# Notes

- A block on storage devices loaded into a **page** in main memory
  - We sometimes interchange page with block
- Buffer pages
  - Often refer to pages in main memory used to store input, output, and intermediate data for an algorithm
- Run: a **sorted sublist** of input data

# Notes


- Make a pass through data:
  - Loading the entire data from disk once

# Sorting

- Illustrates the difference in algorithm design when your data is not in main memory:
  - Problem: sort 1TB of data with 1GB of RAM
- Arises in many places in database systems:
  - Data requested in sorted order (ORDER BY)
  - Needed for grouping operations
  - First step in sort-merge join algorithm
  - Duplicate removal
  - Bulk loading technique for creating B+-tree indexes

# 2-Way Merge-sort: Requires 3 Buffers

- Pass 0: Read a page, sort it, write it
  - only one buffer page is used
- Pass 1, 2, ..., etc.: merging two runs at a time
  - three buffer pages used.



# Two-Way External Merge Sort

- Each pass we read + write each page in file.

- N pages in the file => the number of passes

$$= \lceil \log_2 N \rceil + 1$$

- So total cost is:

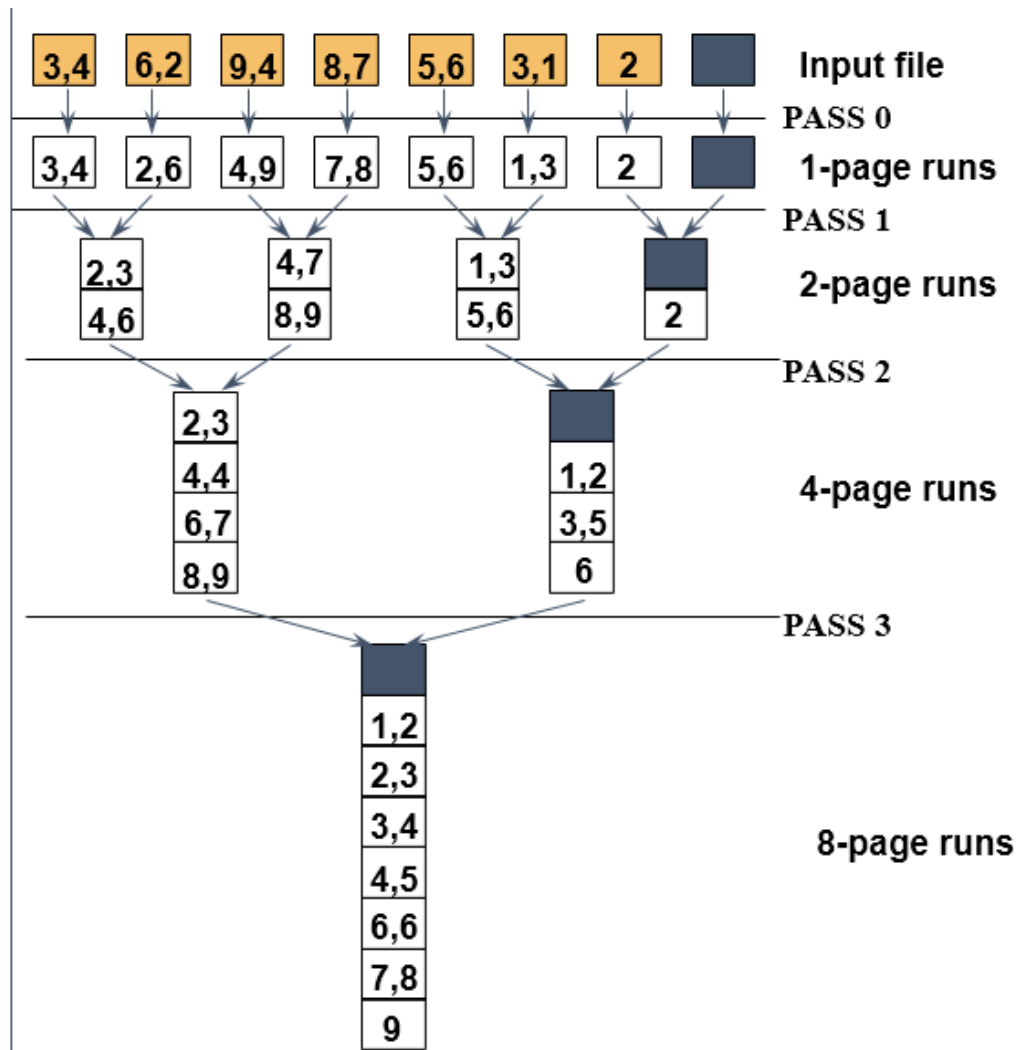
$$2N(\lceil \log_2 N \rceil + 1)$$

- Sort 4MB with buffer page size = 4KB: needs 11 passes

$$N=7, k=3$$

$$1 * 2^k \geq N$$

$$k = \lceil \log_2 (N) \rceil = \lceil 2.8 \rceil = 3$$



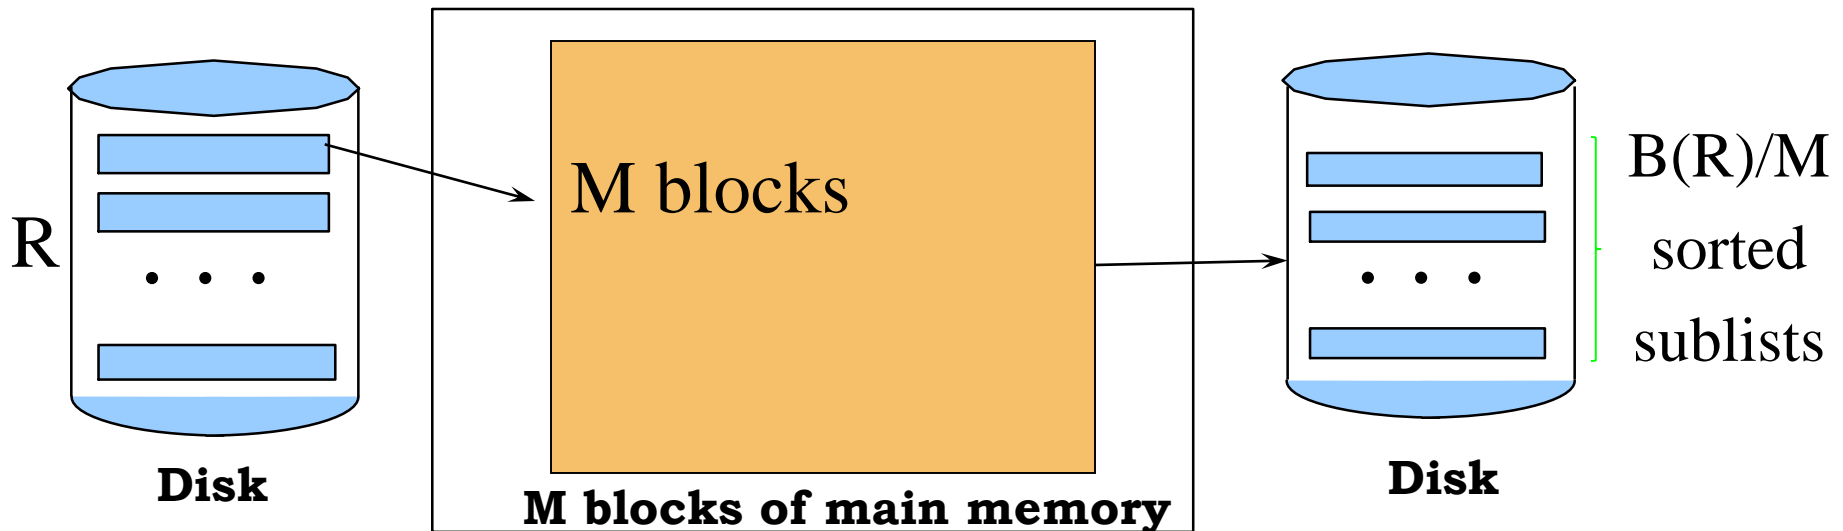
# Can We Do Better ?

- We have more main memory
- Should use it to improve performance
  
- $M$ : # of blocks (i.e., pages) in main memory
- $B(R)$ : # of blocks of relation  $R$



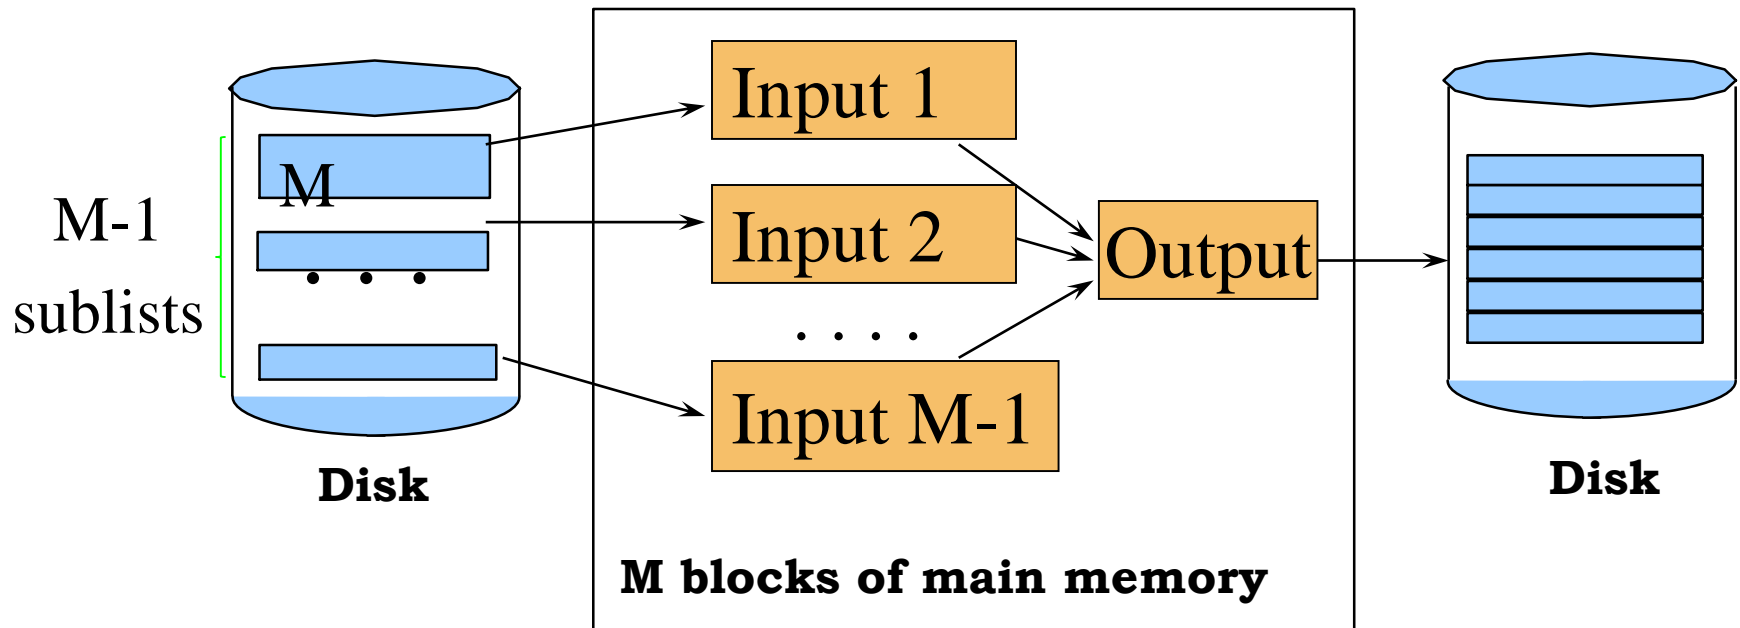
# External Merge-Sort

- Pass 0: load  $M$  blocks in memory, sort
  - Result:  $B(R)/M$  sorted sublists of size  $M$
  - Each sorted sublist is a run



# Pass One

- Merge  $M - 1$  runs into a new run
- Result: each run has now  $M (M - 1)$  blocks

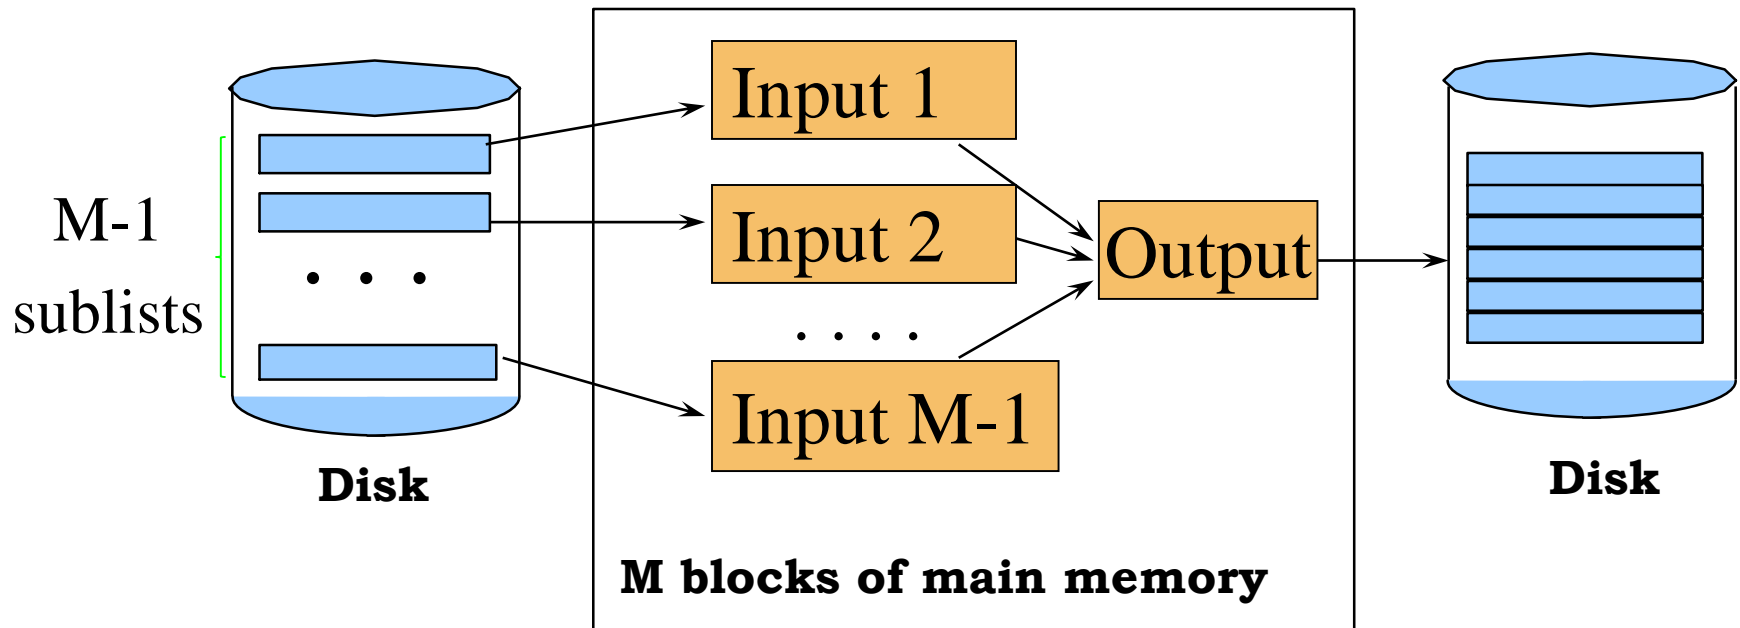


# Cost of Two-Pass, Multiway Merge Sort

- Pass 0: sort  $B/M$  sublists of size  $M$ , write
  - Cost:  $2B(R)$
- Pass 1: merge  $B/M$  sublists, write
  - Cost:  $2B(R)$
- Total cost:  $4B(R)$
- Assumption:  $B(R) \leq M^2$ 
  - $B/M \leq M - 1$  or
  - $B \leq M(M-1) \sim M^2$

# Generalized to k Passes


- Merge every  $M - 1$  runs into a new run
- Result: each run has now  $M (M - 1)^k$  blocks



# If $k$ is the last pass

- Merge  $M - 1$  runs into a single run
- We must have  $M (M - 1)^k \geq B \rightarrow$

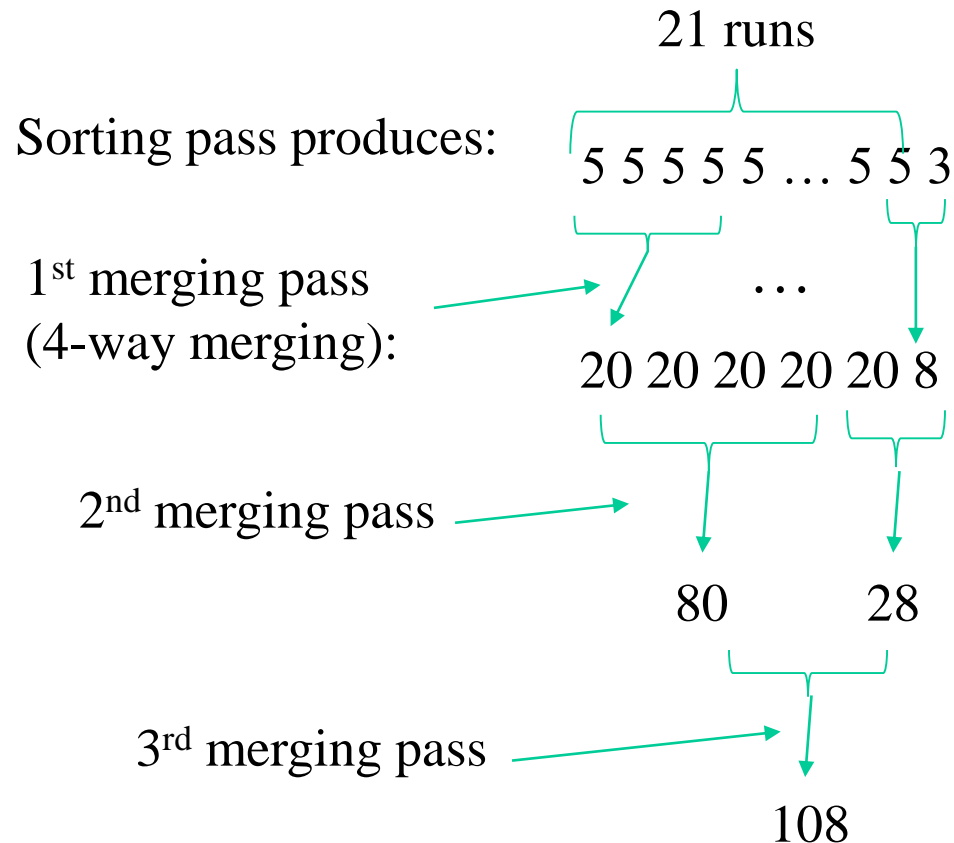
$$k = \lceil \log_{M-1} \lceil B / M \rceil \rceil$$



# Cost of External Merge Sort

- Number of passes:  $1 + \lceil \log_{M-1} \lceil B / M \rceil \rceil$
- Cost =  $2B * (\# \text{ of passes})$
- E.g., with 5 buffer pages, to sort 108 page file:
  - Pass 0: produces  $\lceil 108/5 \rceil = 22$  runs (21 sorted runs of 5 pages each + last run of only 3 pages)
  - Pass 1:  $\lceil 22/4 \rceil = 6$  (5 sorted runs of 20 pages each + last run of only 8 pages)
  - Pass 2: 2 sorted runs, 80 pages and 28 pages
  - Pass 3: Sorted file of 108 pages

# Example Illustrated



# Sorting 1TB using 1GB Memory

- $B(R) = 1\text{TB}/4\text{KB}$  (blocks),  $M = 1\text{GB}/4\text{KB}$  (pages)
- Sorting phase produces  $1024$  runs =  $1\text{K}$  runs
- Merging:
  - Can do:  $1\text{GB}/4\text{KB}-1 = 256\text{K}-1$  ways of merging
  - Can we finish merging in one merging pass?