

```
[ec2-user@ip-172-31-34-154 ~]$ spark-submit q1.py
20/11/03 02:27:38 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

Name
Austria
Belgium
Switzerland
Denmark
Finland
Greece
Netherlands
Norway
Poland
Portugal
Russian Federation
Sweden

```
[ec2-user@ip-172-31-34-154 ~]$ vim q2.py
[ec2-user@ip-172-31-34-154 ~]$ spark-submit q2.py
20/11/03 02:55:17 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

country_name	capital_name
Anguilla	The Valley
Antigua and Barbuda	Saint John's
Aruba	Oranjestad
Bahamas	Nassau
Barbados	Bridgetown
Belize	Belmopan
Bermuda	Hamilton
Canada	Ottawa
Cayman Islands	George Town
Costa Rica	San JosÃ©

```
[ec2-user@ip-172-31-34-154 ~]$ spark-submit q3.py
20/11/03 06:16:14 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

Name	Languages
Anguilla	English
Antigua and Barbuda	English
Aruba	Dutch
Barbados	English
Belize	English
Bermuda	English
Canada	English, French
Cayman Islands	English
Costa Rica	Spanish
Cuba	Spanish

```
[ec2-user@ip-172-31-34-154 ~]$ spark-submit q4.py
20/11/03 04:06:34 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

Continent	avg_le
Europe	75.01
Africa	61.75555555555555
North America	73.85555555555555
South America	71.675
Oceania	78.8
Asia	70.10689655172413

```
[ec2-user@ip-172-31-34-154 ~]$ vim q5.py
[ec2-user@ip-172-31-34-154 ~]$ spark-submit q5.py
20/11/03 04:56:28 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

Continent	cnt
Europe	5
Africa	5
North America	5
Oceania	3

2.1

When facing a range query, it will try to find the first elements that are greater or equal to 10.

At internal nodes, it will find the index the first key that's greater than 10 and then go to the pointer blocks of the same index. If all keys are less than 10, then it will go to the last pointer block.

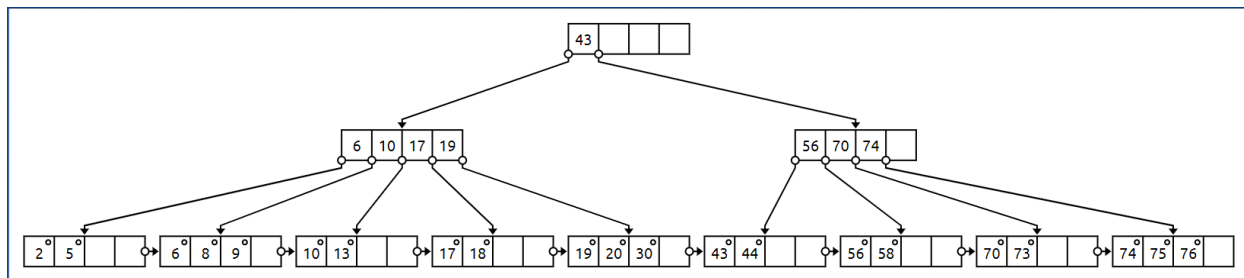
1. First it starts with the root node and finds 43 greater than 10, it will go to the first pointer block.

2. It will parse 6, 10 and find 17 as the first key that's greater than 10. It will then go to the third pointer block.

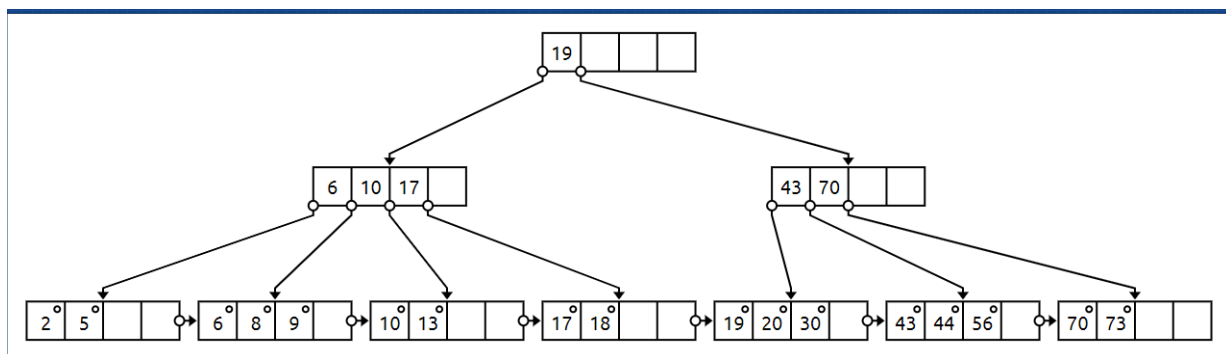
3. When it reaches the leaf level of the tree, it will traverse the linked list from the leftmost pointer of the node and find every number that's ≥ 10 and ≤ 60 . It will stop after seeing a number that's greater than 60. So, it will go through 10, 13, 17, 18, 19, 20, 30, 43, 44, 56, 58 and end up when finding $70 > 60$.

The total number of I/O blocks are 8 blocks.

2.2



2.3



3

- 3.
- (a) Outer loop runs $B(R)/(M-2)$ times and each time will read $B(S)$ once, in total, we read $B(R)$ exactly once. The cost is $B(R) + B(R)/(M-2) * B(S)$
 $= 500 + 500/100 * 1000 = 5500$
 - (b) Sort runs of size M , cost: $2B(R)$, then merge $M-1$ runs, but include each tuple only once, Cost $B(R) + B(S)$. The cost is $3 B(R) + 3 B(S) = 3*500 + 3*1000 = 4500$
 - (c) Step 1: - Hash S into $M - 1$ buckets
Step 2 - Hash R into $M - 1$ buckets
Step 3 - Join every pair of corresponding buckets
The cost is $3 B(R) + 3 B(S) = 4500$
 - (d) Clustered index on attribute a : cost = $B(R)/V(R,a)$, then we are simply replace $B(R)/(M-2)$ in the Nested-loop join with $B(R)/V(R,a)$ and the other steps follow the same. So total cost will be $B(R) + T(R)B(S)/V(S,a) = 500 + 10000*1000/20 = 500500$
- Both b and c are the most efficient ones.